

Quantitative text analysis: Word Embeddings

Blake Miller

MY 459: Quantitative Text Analysis

March 23, 2019

Course website: lse-my459.github.io

1. Overview and Fundamentals
2. Descriptive Statistical Methods for Text Analysis
3. Automated Dictionary Methods
4. Machine Learning for Texts
5. Supervised Scaling Models for Texts
6. *Reading Week*
7. Unsupervised Models for Scaling Texts
8. Similarity and Clustering Methods
9. Topic models
10. Word embeddings
11. Working with Social Media

Overview of text as data methods



Outline

- ▶ Overview of topic models
- ▶ Latent Dirichlet Allocation (LDA)
- ▶ Validating the output of topic models
- ▶ Examples
- ▶ Choosing the number of topics
- ▶ Extensions of LDA

Extensions of LDA

1. Structural topic model (Roberts et al, 2014, AJPS)
2. Dynamic topic model (Blei and Lafferty, 2006, ICML; Quinn et al, 2010, AJPS)
3. Hierarchical topic model (Griffiths and Tenenbaum, 2004, NIPS; Grimmer, 2010, PA)

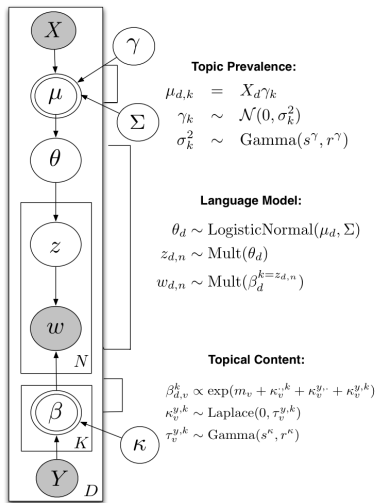
Why?

- ▶ Substantive reasons: incorporate specific elements of DGP into estimation
- ▶ Statistical reasons: structure can lead to better topics.

Structural topic model (STM)

- ▶ Basic idea: $STM = LDA + \text{Contextual Information}$
- ▶ STM provides two ways to include contextual information
 - ▶ Topic prevalence can vary by metadata (e.g. Democrats talk more about education than Republicans)
 - ▶ Topic content can vary by metadata (e.g. Democrats are less likely to use the word “life” when talking about abortion than Republicans)
- ▶ Including context improves the model:
 - ▶ more accurate estimation
 - ▶ better qualitative interpretability

Structural topic model (STM)




- **Prevalence:** Prior on the mixture over topics is now document-specific, and can be a function of covariates (documents with similar covariates will tend to be about the same topics)
- **Content:** distribution over words is now document-specific and can be a function of covariates (documents with similar covariates will tend to use similar words to refer to the same topic)

Structural topic model (STM)


- ▶ User specifies the number of topics: K
- ▶ Observed data for standard topic models
 - ▶ Each document ($d \in 1 \dots D$) is a collection of N_d tokens
 - ▶ Each token is a particular word from a dictionary of V entries
 - ▶ Data summarized in a single matrix $D \times V$ matrix \mathbf{W}
- ▶ Additional data for STM
 - ▶ Topic prevalence covariates: $D \times P$ matrix \mathbf{X}
 - ▶ Topical content groups: D length vector \mathbf{Y}
- ▶ Latent variables
 - ▶ $D \times K$ matrix θ : proportion of document on each topic.
 - ▶ $K \times V$ matrix β : probability of drawing a word conditional on topic.
 - ▶ Low rank approximation to expected counts: $\tilde{W}_{D \times V} \sim \theta_{D \times K} \beta_{K \times V}$

Dynamic topic model




Source: Blei, "Modeling Science"



Dynamic topic model



Source: Blei, “Modeling Science”

Figure 5. Two topics from a dynamic topic model. This model was fit to *Science* from 1880 to 2002. We have illustrated the top words at each decade.





Outline

- ▶ Extensions of LDA
- ▶ Word embeddings:
 - ▶ Overview
 - ▶ Applications
 - ▶ Bias
 - ▶ Embeddings demo
- ▶ Encoding issues

Beyond bag-of-words

Most applications of text analysis rely on a **bag-of-words** representation of documents


- ▶ Only relevant feature: frequency of features
- ▶ Ignores context, grammar, word order...
- ▶ Wrong but often irrelevant

An alternative: Word Vectors


- ▶ One alternative to the bag-of-words approach is **word embeddings**
- ▶ **Word embeddings** represent words as **real-valued vector** in a multidimensional space (often 100–500 dimensions)
- ▶ Central logic:
 - ▶ “You shall know a word by the company it keeps” (Rupert Firth)
 - ▶ Synonyms like oculist and eye-doctor tend to occur in the same environment.
 - ▶ The difference in meaning between two words corresponds “roughly to the amount of difference in their environments” (Harris, 1954, 157).
- ▶ Word vectors learn vector representations of words from the context in which they appear.

Word embeddings example


word	D_1	D_2	D_3	...	D_N
man	0.46	0.67	0.05
woman	0.46	-0.89	-0.08
king	0.79	0.96	0.02
queen	0.80	-0.58	-0.14



Male-Female




Verb tense




Country-Capital

word2vec (Mikolov 2013)

- ▶ Statistical method to efficiently learn word embeddings from a corpus, developed by Google engineer
- ▶ Most popular, in part because pre-trained vectors are available.
- ▶ Two models to learn word embeddings:



CBOW



Skip-gram

How word2vec works

... lemon, a [tablespoon of apricot jam, a] pinch ...
c1 c2 t c3 c4

How word2vec works

- ▶ Main idea: Train a classifier on a binary prediction task:
 - ▶ Is w likely to show up near "apricot"?
- ▶ We don't actually care about this task, but we'll take the learned classifier weights as the word embeddings

How word2vec works

- ▶ A word s near apricot acts as "correct answer" to the question
 - ▶ "Is word w likely to show up near apricot?"
- ▶ No need for hand-labeled supervision
- ▶ The idea comes from neural language modeling
 - ▶ Bengio et al. (2003)
 - ▶ Collobert et al. (2011)

Skip-gram algorithm

1. Treat the target word and a neighboring context word as positive examples.
2. Randomly sample other words in the lexicon to get negative samples
3. Use logistic regression to train a classifier to distinguish those two cases
4. Use the weights as the embeddings

Setup

... lemon, a [tablespoon of apricot jam, a] pinch ...
 c1 c2 t c3 c4

- ▶ Let's represent words as vectors of some length (say 300), randomly initialized.
- ▶ So we start with $300 \times V$ random parameters
- ▶ Over the entire training set, we'd like to adjust those word vectors such that we
 - ▶ Maximize the similarity of the target word, context word pairs (t,c) drawn from the positive data
 - ▶ Minimize the similarity of the (t,c) pairs drawn from the negative data.

Training Data

positive examples +

t	c
apricot	tablespoon
apricot	of
apricot	preserves
apricot	or


negative examples -

t	c	t	c
apricot	aardvark	apricot	twelve
apricot	puddle	apricot	hello
apricot	where	apricot	dear
apricot	coaxial	apricot	forever

How to learn word2vec (skip-gram) embeddings

- ▶ Start with V random 300-dimensional vectors as initial embeddings
- ▶ Use logistic regression, the second most basic classifier used in machine learning after naive bayes
- ▶ Take a corpus and take pairs of words that co-occur as positive examples
- ▶ Take pairs of words that don't co-occur as negative examples
- ▶ Train the classifier to distinguish these by slowly adjusting all the embeddings to improve the classifier performance
- ▶ Throw away the classifier and keep the embeddings.


Vectors and Semantic Relationships



Male-Female





Verb tense



Country-Capital

Vectors capture some general semantic information about words and their relationships to one another.

Analogy: Embeddings capture relational meaning!



- ▶ $\text{vector}(\text{'king'}) - \text{vector}(\text{'man'}) + \text{vector}(\text{'woman'}) \approx \text{vector}(\text{'queen'})$
- ▶ $\text{vector}(\text{'Paris'}) - \text{vector}(\text{'France'}) + \text{vector}(\text{'Italy'}) \approx \text{vector}(\text{'Rome'})$

Outline

- ▶ Extensions of LDA
- ▶ Word embeddings:
 - ▶ Overview
 - ▶ Applications
 - ▶ Bias
 - ▶ Embeddings demo
- ▶ Encoding issues

Application: Pomeroy et al 2018





Figure 4: *Distances by core countries*. Plot of Euclidian distances between US and Russia (gray), and US and China (maroon).

Application: semantic shifts

Using word embeddings to visualize changes in word meaning:




Source: Hamilton et al, 2016 ACL.

<https://nlp.stanford.edu/projects/histwords/>

Application: semantic shifts

Using word embeddings to visualize changes in word meaning:



Source: Hamilton et al, 2016 ACL.

<https://nlp.stanford.edu/projects/histwords/>

Application: dictionary expansion

Using word embeddings to expand **dictionaries** (e.g. incivility)

```
> distance(file_name = "FBvec.bin",  
+         search_word = "libtard",  
+         num = 10)
```

Entered word or sentence: libtard

Word: libtard Position in vocabulary: 5753

	word	dist
1	lib	0.798957586288452
2	lefty	0.771853387355804
3	libturd	0.762575328350067
4	teabagger	0.744283258914948
5	teabilly	0.715277075767517
6	liberal	0.709996342658997
7	retard	0.690707504749298
8	dumbass	0.690422177314758
9	rwnj	0.684058785438538
10	republitard	0.678197801113129

```
> distance(file_name = "FBvec.bin",  
+         search_word = "idiot",  
+         num = 10)
```

Entered word or sentence: idiot

Word: idiot Position in vocabulary: 646

	word	dist
1	imbecile	0.867565214633942
2	asshole	0.848560094833374
3	moron	0.781079053878784
4	asshat	0.772150039672852
5	a-hole	0.765781462192535
6	ahole	0.760824918746948
7	asswipe	0.742586553096771
8	ignoramus	0.735219776630402
9	arsehole	0.732272684574127
10	idoit	0.720151424407959

Outline

- ▶ Extensions of LDA
- ▶ Word embeddings:
 - ▶ Overview
 - ▶ Applications
 - ▶ Bias
 - ▶ Embeddings demo
- ▶ Encoding issues

Bias in word embeddings

Semantic relationships in embeddings space capture stereotypes:

- ▶ Neutral example: man – woman \approx king – queen
- ▶ Biased example: man – woman \approx computer programmer – homemaker

Gender stereotype *she-he* analogies.

sewing-carpentry	register-nurse-physician	housewife-shopkeeper
nurse-surgeon	interior designer-architect	softball-baseball
blond-burly	feminism-conservatism	cosmetics-pharmaceuticals
giggle-chuckle	vocalist-guitarist	petite-lanky
sassy-snappy	diva-superstar	charming-affable
volleyball-football	cupcakes-pizzas	hairstylist-barber

Gender appropriate *she-he* analogies.

queen-king	sister-brother	mother-father
waitress-waiter	ovarian cancer-prostate cancer	convent-monastery

Source: Bolukbasi et al, 2016. arXiv:1607.06520

See also Garg et al, 2018 PNAS and Caliskan et al, 2017 Science.

Outline

- ▶ Extensions of LDA
- ▶ Word embeddings:
 - ▶ Overview
 - ▶ Applications
 - ▶ Bias
 - ▶ Embeddings demo
- ▶ Encoding issues

Outline

- ▶ Extensions of LDA
- ▶ Word embeddings:
 - ▶ Overview
 - ▶ Applications
 - ▶ Bias
 - ▶ Embeddings demo
- ▶ Encoding issues

Character encodings

- ▶ **Encoding:** how digital binary signals are translated into human-readable characters.
 - e.g. 0100100 is displayed as 'd'
- ▶ This also includes characters such as á, ç, ü, etc.
- ▶ **Problem:** many different translation tables, sometimes hard to know which one is used
- ▶ R works with the default encoding scheme in your system:

```
> Sys.getlocale(category = "LC_CTYPE")  
[1] "en_US.UTF-8"
```
- ▶ For English Mac and Linux systems, generally UTF-8. For Windows systems, Windows-1252.
- ▶ UTF-8 (part of Unicode standard) is most popular scheme and used on many websites.